

CS342 Operating Systems – Fall 2018

Project 4: Disk Scheduling

Assigned: Dec 14, 2018

Due date: Dec 29, 2018, **23:55** (Moodle)

You will do this project individually. You have to program in C and Linux. You are recommended to use the following distribution of Linux: **Ubuntu 16.04 – 64 bit**.

Part A: Processes [35 points]

Objective: Practice mass storage, C programming, statistics knowledge.

In this project, you will write a program (diskschedule.c) that implements the following disk scheduling algorithms. a) FCFS; b) SSTF; c) SCAN; d) C-SCAN; e) LOOK; f) C-LOOK.

Your program will service a disk with 5000 cylinders numbered 0 to 4999. It will service 1000 requests according to each of the algorithms listed above. The program will be passed the initial position of the disk head (as a parameter on the command line) and report the total amount of head movement required by each algorithm. The workload (requests) will be given in two ways: 1) will be randomly generated in your program; 2) will be read from an input file. Which one to use will be specified at the command line. The program will be invoked as follows:

diskschedule <headpos> <inputfile>

The name of the program will be **diskschedule**. <headpos> is the initial head position. Assume the initial direction (when required) is always towards right (to bigger cylinder numbers). If <inputfile> is not given, then the requests will be generated randomly. The format of the input file (ascii text file) is given in the example below. Each line is for a different request. A line contains a request number and a cylinder number.

```
1 4000
2 2000
3 4500
4 1450
5 6534
...
999 452
```

An example invocation of the program can be:

```
diskschedule 1230 in.txt
```

The output file be in the following example format:

```
FCFS: 15000
SSTF: 15000
SCAN: 15000
```

C-SCAN: 15000
LOOK: 15000
C-LOOK: 15000

Part C: Experiments [20 points]

Objective: *Practice designing and conducting experiments and applying knowledge and skills acquired in the Probability and Statistics course.*

Run the program 100 different times with random input (random requests and random initial head position). At the end, for each algorithm, find out the *average* total movement and *standard deviation* of the total movement. Report the results in a table.

Submission

Put all your files into a project directory named with your ID (one of the IDs of team members), tar the directory (using **tar xvf**), zip it (using **gzip**) and upload it to Moodle. For example, a student with ID 20140013 will create a directory named 20140013, will put the files there, tar and gzip the directory and upload the file. The uploaded file will be 20140013.tar.gz. Include a **README.txt** file as part of your upload. It will have the name and ID of the student, at least. Include also a **Makefile** to compile your program. We want to type just **make** and obtain the executables. Do not forget to put your report (PDF form) into your project directory.

Additional Information and Clarifications

- Additional clarifications can be posted in piazza or on the course website besides project specification.